

A PRACTICAL HISTORY OF R (WHERE THINGS CAME FROM)

Roger Bivand

16 May 2018

The keynote presentation files are on github at: [rsbivand/eRum18](#)

- Not infrequently, we wonder why choices such as `stringsAsFactors=TRUE` or `drop=TRUE` were made.
- Understanding the original uses of S and R (in the 1980s and 1990s), and seeing how these uses affected the development of R lets us appreciate the robustness of R's ecosystem.
- This keynote uses readings of the R sources and other information to explore R's history. The topics to be touched on include the “colour” books (brown, blue, white, green), interlinkages to SICP (Scheme) and LispStat
- We'll also touch on the lives of R-core, the mailing lists and CRAN, and Ancients and Moderns (see [Exploring the CRAN social network](#)).

History of R and its data structures

- [Rasmus Bååth](#) has a useful blog piece on R's antecedents in the S language
- Something similar is present in the second chapter of (Chambers, 2016), from the viewpoint of one of those responsible for the development of the S language
- In addition to S, we need to take [SICP and Scheme](#) into account (Abelson and Sussman, 1996, second edition), as described by Ihaka and Gentleman (1996) and Wickham (2014)
- Finally, LispStat and its creators have played and continue to play a major role in developing R (Tierney, 1990, 1996, 2005)


R did not always look like an alternative implementation of the S language. It started as a small Scheme-like interpreter (loosely based on work by Sam Kamin [4] and David Betz [2]). This provided a platform for experimentation and extension. The following dialog shows a very early version of the R interpreter at work.

```
> (define square (lambda (x) (* x x)))
square
> (define v 10)
v
> (square v)
100
```


The S-like appearance of R was added incrementally. We initially moved to an S-like syntax for our experiments because, although we were both familiar with Lisp syntax, we didn't feel that it was really suitable for expressing statistical computations. This choice of syntax set us on a path towards compatibility with S. Once initiated, the move towards compatibility with S was irresistible. This was partly because we wanted to see just how far we could push it and partly because it gave us access to code resources and developers.

Becker and Chambers (1984): S: An Interactive Environment for Data Analysis and Graphics, A.K.A. the Brown Book

Becker and Chambers (1985):
Extending the S System




Becker et al. (1988): The New S Language: A Programming Environment for Data Analysis and Graphics, A.K.A. the Blue Book.
Chambers and Hastie (1992): Statistical Models in S, A.K.A. the White Book.



Chambers (1998): Programming with Data: A Guide to the S Language, A.K.A. the Green Book.

Venables and Ripley (2000): S Programming



- The S2 system was described in the Brown Book, S3 in the Blue Book and completed in the White Book, finally S4 in the Green Book
- The big advance from S2 to S3 was that users could write functions; that data.frame objects were defined; that formula objects were defined; and that S3 classes and method dispatch appeared
- S4 brought connections and formal S4 classes, the latter seen in R in the **methods** package (still controversial)
- S-PLUS was/is the commercial implementation of S and its releases drove S3 and S4 changes

- S was a Bell Labs innovation, like Unix, C, C++, and many interpreted languages (like AWK); many of these share key understandings
- Now owned by Nokia, previously Alcatel-Lucent, Lucent, and AT&T
- Why would a telecoms major (AT&T) pay for fundamental research in computer science and data analysis (not to sell or market other products better)?
- Some Green Book examples are for quality control of telecoms components

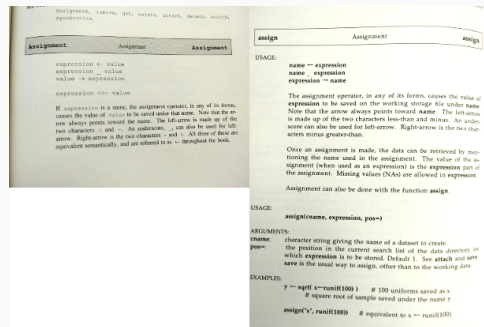
- S-PLUS was quickly adopted for teaching and research, and with S3, provided extensibility in the form of libraries
- Most links have died by now, but see this [FAQ](#) for a flavour - there was a lively community of applied statisticians during the 1990s
- S built on a long tradition of documentation through examples, with use cases and data sets taken from the applied statistical literature; this let users compare output with methods descriptions
- ... so we get to R

- Luke Tierney was in R core in 1997, and has continued to exert clear influence over development
- Because R uses a Scheme engine, similar to Lisp, under the hood, his insight into issues like the garbage collector, namespaces, byte-compilation, serialization, parallelization, and now **ALTREP** has been crucial ([see also the proposal by Luke Tierney, Gabe Becker and Tomas Kalibera](#))
- Many of these issues involve the defensive copy on possible change policy involved in lazy evaluation, which may lead to multiple redundant copies of data being present in memory
- Luke Tierney and Brian Ripley have fought hard to let R load fast, something that is crucial to ease the use of R on multicore systems or inside databases

Vintage R

USE QUESTIONS: LEFT ASSIGN

Why was `underscore_separated` not a permitted naming convention in R earlier (see [Bååth \(2012\)](#))? `_` was not a permitted character in names until it had lost its left assign role, the same as `<-`, in 1.9.0 in 2004. (Brown Book p. 256, Blue Book p. 387)



Why is the **factor** storage mode still so central? **stringsAsFactors = TRUE** was the legacy **as.is = FALSE**; analysis of categorical variables was more important, and **factor** only needed to store **nlevels()** strings (White Book p. 55-56, 567)

If this is on a file named "auto1" it can be read in and turned into a data frame by a simple call to `read.table()`.

```
> somedata <- read.table("auto1")
> dsn(somedata)
[1] 10 5
> dinnames(somedata)
[[1]]:
[1] "Acura.Integra.4"      "Audi.100.5"          "BMW.325i.6"
[4] "Chevrolet.Lumina.4"   "Ford.Festiva.4"       "Mazda.929.V6"
[7] "Mazda.MX.5.Miata"     "Nissan.300ZX.V6"      "Oldsmobile.Calais.4"
[10] "Toyota.Cressida.6"

[[2]]:
[1] "Price"      "Country"      "Reliability" "Mileage"      "Type"
```

Let's look at this example in a little more detail to see what is going on. The file contained fields separated by "white space", one or more blanks or tabs. The first line of the file contained five fields, meant to be the names for the variables; the remaining lines had six fields, the first being the row label and the rest the data for this observation. Some of the fields are numeric; others are character strings. The character strings will be turned into factor variables, as we can see by the following:

56

CHAPTER 3. DATA FOR MODELS

```
> sapply(somedata, data.class)
      Price Country Reliability Mileage      Type
"numeric" "factor" "numeric"  "numeric" "factor"
```


drop = TRUE for array-like objects;
since matrices are vectors with a **dim**
attribute, choosing (part of) a row or
column made **dim** redundant (Blue
Book p. 128, White Book p. 64)

creates a matrix with two rows, since the specified names occur among the *rownames* for the matrix. As with vector subscripts, it is possible to replicate rows and/or columns in subscripting matrices. In fact, this is a powerful technique (see the exercises below).

If the matrices resulting from subscripting end up with one row or one column, there is a choice to make. *S* can retain the matrix properties or it can drop one dimension to produce a vector. By default, *S* drops these redundant, or “dead,” dimensions.

```
> m <- matrix(1:12, 3, 4)
> m[,2]
[1] 4 5 6
```

In most cases of interactive computing, that is the right choice since it tends to reduce complexity—why have a 1-column matrix when a vector will do. However, when you are writing functions, you may need to rely on a subscripted matrix remaining a matrix. In this case, add `drop=F` to the subscripts:

```
m[,which,drop=F]
```

This ensures that no dimensions will be dropped, and allows you to avoid surprises when one or the other subscript happens to leave only one element.

Exercises

Exercises ———
Data frames can be treated as matrices in calls to most of the basic functions treating arrays: subsets and elements, `dim()`, `dimnames()`, and functions based on those. If *x* is a data frame, then

```
x[i,]; x[,j]; x[i,j]
dim(x); dimnames(x)
nrow(x); ncol(x)
```

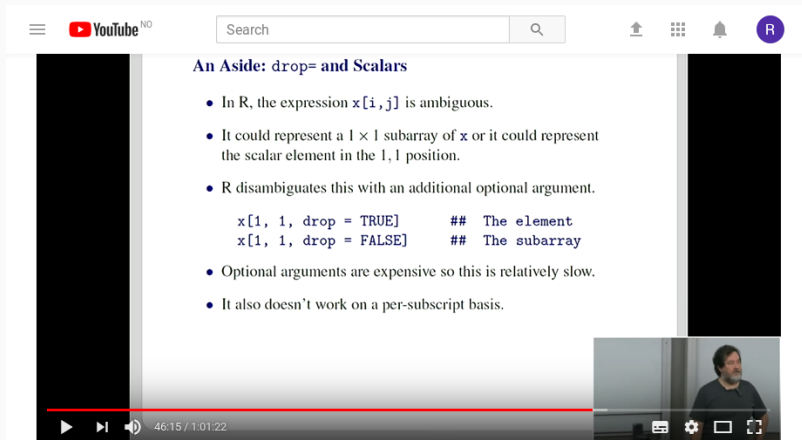
produce results corresponding intuitively to their behavior on matrices. For example, `x[i,]` produces a new data frame by using *i* to index the rows of *x*. The indexing by *i* can use *numeric*, *logical*, or *character* values. Similarly, `x[,j]` indexes on the columns (the variables) and `x[i,j]` on both. When a single column is selected, the result is by default the variable, not a data frame containing one variable. This action can be suppressed by including the argument `drop=F`, following the rules applied to arrays (S, page 128). For example, if *stats* was some statistic computed for each of the variables in *market.frame*, the expression

```
market.frame[, stats > cutoff , drop=F]
```

ensures that the extracted object is still a data frame, even if it has only one column. A single row by default remains a data frame—there is no generally useful object corresponding to rows of a data frame. If you really want to, however, you can cause the single row to be dropped to a list by including the argument `drop=T`.

BUT SCALARS ARE ALSO VECTORS ...

Treating scalars as vectors is not efficient:



The image shows a YouTube video player interface. At the top, there is a search bar with the text "Search" and a magnifying glass icon. To the right of the search bar are icons for upload, grid, notifications, and a user profile icon labeled "R". Below the search bar, the video content is displayed. The title of the video is "An Aside: drop= and Scalars". The video content includes a list of bullet points and two lines of R code. The first bullet point states that in R, the expression `x[i, j]` is ambiguous. The second bullet point explains that it could represent a 1×1 subarray of `x` or the scalar element at the $1, 1$ position. The third bullet point notes that R disambiguates this with an additional optional argument. The code examples show `x[1, 1, drop = TRUE]` for the element and `x[1, 1, drop = FALSE]` for the subarray. The fourth bullet point mentions that optional arguments are expensive, making this relatively slow. The fifth bullet point states that this approach doesn't work on a per-subscript basis. At the bottom of the video player, there is a progress bar showing the video is at 46:15 out of 1:01:22. A small inset video in the bottom right corner shows a man speaking.

YouTube

Search

An Aside: drop= and Scalars

- In R, the expression `x[i, j]` is ambiguous.
- It could represent a 1×1 subarray of `x` or it could represent the scalar element in the $1, 1$ position.
- R disambiguates this with an additional optional argument.

```
x[1, 1, drop = TRUE]    ## The element
x[1, 1, drop = FALSE]   ## The subarray
```

- Optional arguments are expensive so this is relatively slow.
- It also doesn't work on a per-subscript basis.


46:15 / 1:01:22

- An R-0.49 source tarball is available from CRAN
- Diffs for Fedora 27 (gcc 7.3.1) include setting compilers and `-fPIC` in `config.site`, putting `./` before `config.site` in `configure`, and three corrections in `src/unix`: in `dataentry.h` add `#include <X11/Xfuncproto.h>` and comment out `NeedFunctionPrototypes`; in `rotated.c` comment out `/*static*/ double round`; in `system.c` comment out `__setfpucw` twice; BLAS must be provided externally
- Not (yet) working: prototypes are missing in the `eda` and `mva` packages so the shared objects fail to build
- R-0.49 video


R SVN logs

- The command: `svn log --xml --verbose -r 6:74688 https://svn.r-project.org/R/trunk > trunk_verbose_log1.xml` provides a rich data source
- Each log entry has a revision number, author and timestamp, message and paths to files indicating the action undertaken for each file
- The XML version is somewhat easier to untangle than the plain-text version
- I haven't tried possible similar approaches to Winston Chang's [github r-source repo](#)

COMMITTS 1998-2017



COMMITTS BY AUTHOR AND YEAR BETWEEN R6 AND R74688



```
## <logentry revision="6">
##   <author>ihaka</author>
##   <date>1997-09-18T04:41:25.000000Z</date>
##   <paths>
##     <path action="M" prop-mods="false" text-mods="true" kind="file">/t
##   </paths>
##   <msg>New predict.lm from Peter Dalgaard</msg>
## </logentry>
```



```
## <logentry revision="74688">
##   <author>ripley</author>
##   <date>2018-05-03T11:11:00.501520Z</date>
##   <paths>
##     <path text-mods="true" kind="file" action="M" prop-mods="false">/t
##   </paths>
##   <msg>TexLive 2018 may produce logs not in the current encoding</msg>
## </logentry>
```

COMMIT MESSAGES BY NUMBER OF FILES AFFECTED, YEAR AND REVISION

```
## 2015 68948 2150 use https
## 2012 59039 1727 use preferred form of 'R Core Team'
## 2011 56186 1260 Revert r56184 and r56185
## 2011 56184 1249 Remove redundant \alias entries from man pages
## 2007 42333 1223 add copyright/licence header, remove CVS-style $Id f
## 2012 61433 620 remove trailing spaces
## 2012 60146 602 add copyright statements
## 2007 42338 559 add licence statements
## 2012 59780 524 update, including bug-reporting address
## 2003 27444 497 splitting base
```

##				
##	tools	FAQ	m4	etc
##	396	467	579	601
##	configure.ac	share	configure	BUGS
##	693	1085	1319	1485
##	date-stamp	po	NEWS	tests
##	2531	3762	5842	11445
##	doc	src		
##	11484	97067		


##						
##	windows	graphics	gnome	macintosh	appl	unix
##	74	79	217	611	796	1524
##	nmath	scripts	extra	modules	include	gnuwin32
##	1764	1966	2315	2358	3528	9140
##	main	library				
##	15011	57563				

EPOCH FILE COMMITS IN TRUNK/SRC/LIBRARY

##				
##	compiler	profile	stats4	translations
##	234	247	261	296
##	nls	datasets	modreg	mva
##	319	333	402	462
##	splines	ctest	tcltk	ts
##	464	549	865	900
##	parallel	grid	graphics	grDevices
##	1096	1931	2158	3695
##	methods	utils	stats	tools
##	3982	5397	6840	7028
##	base			
##	19331			

##				
##	Makefile	DESCRIPTION.in	makebasedb.R	Makefile.win
##	8	10	11	18
##	baseloader.R	demo	data	Makefile.in
##	32	61	66	75
##	inst	po	R	man
##	270	497	6719	11557

FILES BY YEAR AND COMMIT ACTION



CRAN and Bioconductor packages

- Once S3 permitted extension by writing functions, and packaging functions in libraries, S and R ceased to be monolithic
- In R, a library is where packages are kept, distinguishing between base and recommended packages distributed with R, and contributed packages
- Contributed packages can be installed from CRAN (infrastructure built on CPAN and CTAN for Perl and Tex), Bioconductor, other package repositories, and other sources such as github
- With over 12000 contributed packages, CRAN is central to the R community, but is stressed by dependency issues (CRAN is not run by R core)

- Andrie de Vries [Finding clusters of CRAN packages using igraph](#) looked at CRAN package clusters from a page rank graph
- We are over three years further on now, so updating may be informative
- However, this is only CRAN, and there is the big Bioconductor repository to consider too
- Adding in the Bioconductor (S4, curated) repo does alter the optics, as you'll see, over and above the cluster dominated by **Rcpp**

##	Rcpp	MASS	ggplot2	AnnotationDbi
##	0.022891	0.012087	0.011861	0.011142
##	Matrix	dplyr	plyr	mvtnorm
##	0.006958	0.006744	0.005359	0.005113
##	Biobase	survival	stringr	data.table
##	0.005049	0.005023	0.004682	0.004223
##	lattice	igraph	RcppArmadillo	httr
##	0.004112	0.003993	0.003944	0.003898
##	magrittr	jsonlite	IRanges	reshape2
##	0.003896	0.003831	0.003661	0.003515

##	MASS	ggplot2	plyr	mvtnorm
##	0.012086924	0.011860517	0.005359345	0.005113266
##	survival	lattice	igraph	reshape2
##	0.005022831	0.004111811	0.003992951	0.003515287
##	foreach	sp	RColorBrewer	doParallel
##	0.003309872	0.003280003	0.003092672	0.002297795

##	dplyr	stringr	httr	magrittr	jsonlite
##	0.006744351	0.004681562	0.003898281	0.003895957	0.003830517
##	shiny	tidyr	tibble	XML	RCurl
##	0.002985308	0.002493951	0.002476508	0.002418055	0.002164565
##	purrr	lubridate			
##	0.001852909	0.001549448			

THIRD PACKAGE CLUSTER

##	Biobase	IRanges	S4Vectors
##	0.005048914	0.003661421	0.003473939
##	GenomicRanges	Biostrings	BiocGenerics
##	0.002733640	0.002583678	0.002437776
##	DBI	RSQLite	limma
##	0.002257874	0.001849297	0.001582772
##	GenomeInfoDb	SummarizedExperiment	oligoClasses
##	0.001459829	0.001450142	0.001209420

##	Rcpp	Matrix	data.table	RcppArmadillo
##	0.0228907439	0.0069578788	0.0042227075	0.0039442804
##	RcppEigen	BH	rstan	RcppParallel
##	0.0011843649	0.0011622371	0.0004800493	0.0004186803
##	bigmemory	Rdpack	magic	RSpectra
##	0.0003684402	0.0003410619	0.0002661583	0.0002060934

##	AnnotationDbi	org.Hs.eg.db	GenomicFeatures	org.Mm.eg.db
##	1.114150e-02	1.233367e-03	1.098628e-03	6.799030e-04
##	org.Rn.eg.db	geneXtender	ChIPQC	rCGH
##	4.536656e-04	2.306122e-04	1.574986e-04	1.357878e-04
##	chimera	Homo.sapiens	OrganismDbi	TxDb.Hsapiens.U
##	1.335513e-04	1.280833e-04	1.186413e-04	9.027467e-05


##	Formula	sandwich	lmtest	texreg
##	0.0010315610	0.0006274957	0.0005497347	0.0003765079
##	maxLik	prediction	stargazer	partykit
##	0.0003358931	0.0003106929	0.0002972425	0.0002900618
##	penalized	miscTools	AER	plm
##	0.0002295456	0.0002291684	0.0002026993	0.0001901027

A word cloud of R packages. The most prominent package is MASS, followed by ggplot2. Other visible packages include survival, reshape2, doParallel, foreach, igraph, codal, lattice, zoo, scales, numDeriv, gridExtra, nlme, sp, gplots, raster, plyr, mvtnorm, and RColorBrewer.

MASS
reshape2
doParallel
foreach
survival
igraph
codal
lattice
zoo
scales
numDeriv
gridExtra
nlme
sp
gplots
raster
plyr
ggplot2
mvtnorm
RColorBrewer

A word cloud of R packages. The most prominent package is stringr, followed by dplyr. Other visible packages include lubridate, httr, XML, curl, tidy, shiny, R6, jsonlite, RCurl, magrittr, assertthat, knitr, readr, rlang, digest, purrr, and tibble.

stringr
lubridate
httr
XML
curl
tidy
shiny
R6
jsonlite
RCurl
magrittr
assertthat
knitr
readr
rlang
digest
purrr
tibble
dplyr




AnnotationDbi


Formula

- Francois Keck explored CRAN package co-authorship in a more recent blog: [Exploring the CRAN social network](#)
- Once again, a little time has passed, so maybe things have shifted
- Thanks to Martin Morgan, I've added listings corresponding in part to `tools::CRAN_package_db()`
- It is refreshing to see that Bioconductor is clearly present, and the people implicated are active in upgrading R internals

FIRST TWO PACKAGE AUTHOR CLUSTERS

##	Name	Package
## 1	Kurt Hornik	62
## 2	Martin Maechler	52
## 3	Achim Zeileis	49
## 4	Dirk Eddelbuettel	48
## 5	Brian Ripley	35
## 6	Romain Francois	28
## 7	Torsten Hothorn	28
## 8	Ben Bolker	27
## 9	Douglas Bates	26
## 10	Roger Bivand	25
## 11	Thomas Lumley	24
## 12	Michael Friendly	21

##	Name	Package
## 1	Rstudio	110
## 2	Hadley Wickham	107
## 3	Inc	50
## 4	Scott Chamberlain	49
## 5	Jeroen Ooms	40
## 6	R. Core Team	38
## 7	Yihui Xie	36
## 8	Bob Rudis	32
## 9	Jj Allaire	31
## 10	Kirill Muller	31
## 11	Gabor Csardi	30
## 12	Winston Chang	27


THIRD AND FOURTH

##	Name	Package	##	Name	Package
## 1	Bioconductor	Package 37	## 1	Wolfgang Huber	32
## 2	Marc Carlson	29	## 2	Robert Gentleman	28
## 3	Martin Morgan	28	## 3	Rafael A. Irizarry	28
## 4	Herve Pages	22	## 4	Kasper Daniel Hansen	20
## 5	Bioconductor Core Te	19	## 5	Matthew N. McCall	20
## 6	Seth Falcon	15	## 6	Hector Corrada Bravo	15
## 7	Lihua Julie Zhu	13	## 7	Andy Lynch	14
## 8	Aaron Lun	12	## 8	Mark Dunning	13
## 9	Pierre Neuvial	11	## 9	John D. Storey	12
## 10	Gordon Smyth	10	## 10	Andrew E. Jaffe	11
## 11	Valerie Obenchain	10	## 11	Jeffrey T. Leek	11
## 12	Jianhong Ou	9	## 12	Matthew Eldridge	10

FIFTH AND SIXTH


##	Name	Package	##	Name	Package
## 1	Benjamin Haibe-Kains	21	## 1	R. Gentleman	18
## 2	Gianluca Bontempi	18	## 2	Raphael Gottardo	14
## 3	John Quackenbush	12	## 3	Arnaud Droit	13
## 4	Levi Waldron	12	## 4	Greg Finak	10
## 5	Marcel Ramos	10	## 5	H. Pages	7
## 6	Aedin Culhane	9	## 6	Mike Jiang	7
## 7	Catharina Olsen	9	## 7	Astrid Deschenes	6
## 8	Markus Schroeder	8	## 8	Pascal Belleau	6
## 9	Christos Sotiriou	7	## 9	F. Hahne	6
## 10	Ludwig Geistlinger	7	## 10	S. Falcon	6
## 11	Deena M.A. Gendoo	6	## 11	Charles Joly Beaupar	5
## 12	Houtan Noushmehr	5	## 12	N. Gopalakrishnan	5

FIRST TWO PACKAGE AUTHOR CLUSTERS



THIRD AND FOURTH PACKAGE AUTHOR CLUSTERS

Bioconductor Package Maintainer



FIFTH AND SIXTH PACKAGE AUTHOR CLUSTERS

Gianluca Bontempi
Benjamin Haibe-Kains
Houtan Nourshahr
John Quackenbush
Ludwig Geistlinger
Levi Waldron
Rafael A. Irizarry
Aedin Culhane
Markus Schroeder
Lucas Schiffer
Markus Kister
Christos Sotiropoulos
Deena M.A. Gendoo
Catharina Olsen
Marcel Ramos

F. Hahn
Charles Joly Beauparlant
R. Gentleman Greg Finak
Pascal Belleau Arnaud Droit
Raphael Gottardo
N. Gopalakrishnan
Mike Jiang H. Pages
P. Abouyoun
S. Falcon
N. Le Meur
Astrid Deschenes

- Many sources in applied statistics with an S-like syntax but Lisp/Scheme-like internals, and sustained tensions between these
- Many different opinions on preferred ways of structuring data and data handling, opening for adaptations to different settings
- More recently larger commercial interest in handling large input long data sets, previously also present; simulations also generate large output data sets; bioinformatics both wide and long
- Differing views of the world in terms of goals and approaches
- Differences provide ecological robustness